

Rishith Prathi

774-393-1667 | rishiprathi18@gmail.com | [linkedin.com/in/rishith-prathi](https://www.linkedin.com/in/rishith-prathi) | github.com/Rishith18

EDUCATION

Carnegie Mellon University

Pittsburgh, PA

B.S. in Computer Science / Concentration in Machine Learning

Expected May 2028

- Coursework: 15-259 Probability and Computing, 15-213 Computer Systems, 15-251 Discrete Math, 15-210 Parallel and Sequential Data Structures and Algorithms, 10-301 Machine Learning, 10-714 Deep Learning Systems

EXPERIENCE

Adobe

May 2026 – Aug. 2026

Software Engineering Intern

San Jose, CA

- Shipped a production **RAG** context engine with containerized **FastAPI** serving **1,000+** users, generating AI recommendations for **20k** marketing campaigns via Azure OpenAI embeddings and Cosmos DB vector search
- Created a **Scala** service, scheduled through **Jenkins**, that leverages an LLM to cluster AI-agent conversations from Langfuse into weekly Slack reports, saving **10+** engineers a combined **20+** hours/week of manual review
- Migrated database syncing method to utilize distributed **Temporal** workflows with checkpointing and automatic retries, making syncs resumable after failure and eliminating **5** minutes of redundant work per API deploy
- Built an evaluation harness comparing **12+** campaign embedding strategies with LLM-as-judge scoring; the top strategy improved average recommendation accuracy by **27%** and surveyed user satisfaction by **8%**

NeuLab @ Language Technologies Institute/CMU

Jan. 2026 – Jul. 2026

Machine Learning Research Assistant

Pittsburgh, PA

- Co-designed **Forger** (2nd author of paper, COLM 2026), an open-source framework that provides reusable workflows (using web data retrieval + LLM orchestration) for generating LLM training data at scale
- Engineered a thread-safe caching layer sharing GPU-hosted LLMs (tensor-parallel) across concurrent evaluation workers, eliminating redundant weight reloads and cutting eval runtime **3x**
- Demonstrated Forger-generated data drives a **24-point** in-context-learning gain on HumanEval LLM benchmark through batched vLLM inference with seeded 4-shot subset sampling and bootstrapped confidence intervals

Moss Robotics (Seed-Stage Startup)

Jul. 2025 – Dec. 2025

Software Engineering Intern

Pittsburgh, PA

- Programmed a point cloud **training data generator** to improve company's primary ML model accuracy by **4%**
- Boosted YOLO-based **neural network** performance by **23%** by building a C++/Python pipeline that generated and rasterized **5,000+** synthetic images for YOLO object detection
- Validated synthetic data at **89%** of real-data accuracy, reducing **100+** hours of costly real-world data collection

PROJECTS

Phantom - VentureHacks Invite-Only 2026 Winner (Top 2 of 50+)

Mar. 2026

- Built Phantom, an **AI agent** for Meta Ray-Ban smart glasses that handles everyday tasks via voice and vision
- Streamed live camera + audio to **Gemini Live API** via websockets for sub-second multimodal inference
- Engineered an LLM tool-calling pipeline through **OpenClaw**, routing Gemini function calls to native iOS actions, cutting voice-to-action latency by **around 40%** for performance optimization

Splitpot - Distributed CFR Poker Engine

Feb. 2026 – Mar. 2026

- Parallelized a Counterfactual Regret Minimization poker engine in Python across **66K+** information sets on a Kubernetes cluster of autoscaling worker pods, cutting model training time from hours to minutes
- Built a containerized strategy storage system that aggregates results from parallel workers with checkpointing

SwiftER - NexHacks 2026 Winner (Top 3 of 1,000+)

Jan. 2026

- Developed a **multi-agent AI pipeline** with LangChain and GPT-4 that transcribes live EMT voice calls and auto-selects optimal hospitals for STEMI, stroke, and trauma cases in real time
- Designed the routing platform with React, FastAPI, PostgreSQL, and Redis for **sub-second** hospital matching

TECHNICAL SKILLS

Languages: Python, Java, Scala, TypeScript, JavaScript, C, C++, SML, SQL

Tools: Git, Docker, Kubernetes, Linux, Node.js, Temporal, AWS (S3, Lambda), Azure (OpenAI, Cosmos DB)

Frameworks/Libraries: FastAPI, Flask, GraphQL, React, Next.js

AI/ML: PyTorch, TensorFlow, NumPy, Pandas, Hugging Face Transformers, LangChain, vLLM